

Rain and City Infrastructure: Analyzing NYC 311 Service Requests and Weather Data

Group 4:

Adrian Garcia – adrian.garcia1@baruchmail.cuny.edu

Neha Machamada – neha.machamada@baruchmail.cuny.edu

Adnan Quayyum – adnan.quayyum@baruchmail.cuny.edu

Roland Thomas – roland.thomas@baruchmail.cuny.edu

Zicklin School of Business, Baruch College, City University of New York

CIS4400: Data Warehousing for Analytics

Section: CMWA [44276]

Professor Richard Holowczak

Abstract

This project analyzes water-related 311 complaints in New York City, such as sewer issues, water system problems, water conservation concerns, and storm/flooding complaints, in conjunction with daily rainfall and snowfall data for each of the five boroughs. The goal is to identify patterns and correlations between precipitation events and strain on city infrastructure, and to understand how service responses vary geographically.

We will integrate the 311 Service Requests dataset from NYC Open Data with borough-level weather datasets from Open-Meteo API, transform them into separate data marts, and then combine them into a centralized data warehouse. Analyses will focus on identifying communities that are at higher risk for water-related issues during heavy rain or snow, and evaluating response times by borough and type of complaint. These insights may support data-driven decisions about resource allocation and infrastructure planning to reduce vulnerabilities during severe weather events.

The Problem Statement

NYC experiences extreme weather throughout the year, from heat waves in summer to snowstorms in winter. Heavy rain and snow can overwhelm city infrastructure, leading to flooding, sewer backups, and service interruptions. A major concern is how water-related issues affect neighborhoods differently across boroughs and whether the city responds consistently across geographic areas. By focusing on four categories of water-related 311 complaints: Sewer, Water System, Water Conservation, and Storm/Flooding, and linking them with borough-level daily weather data, we can examine how precipitation correlates with the volume and severity of complaints, as well as with response times. Key performance indicators (KPIs) will measure complaint volumes, rainfall and snowfall levels, geographic clustering of incidents, and average service response times during normal vs. extreme weather days. This analysis will allow us to identify boroughs and communities most affected by water-related issues and suggest where the city may need to improve its preparedness and responsiveness.

Part I: Data Sources

The Main Dataset

[NYC Open Data – 311 Service Requests \(2010 to Present\)](#)

- Rows: A total of about 41 million rows, where each row is a 311 service request.
- Columns: A total of 41 columns, including a unique identifier for each service request.
- Columns of Interest or KPIs: Complaint Type, Incident ZIP, Borough, Longitude, Latitude, Descriptor, Created Date, Closed Date, Resolution Description.
- Filter: Complaint Types = Sewer, Water System, Water Conservation, Storm/Street Flooding

NYC Open Data 311 Service Requests Columns (Excluding Unique ID):

Created Date	Closed Date	Agency	Agency Name
Complaint Type	Descriptor	Location Type	Incident Zip
Incident Address	Street Name	Cross Street 1	Cross Street 2
Intersection Street 1	Intersection Street 2	Address Type	City
Landmark	Facility Type	Status	Due Date
Resolution Description	Resolution Action Updated Date	Community Board	BBL
Borough	X Coordinate (State Plane)	Y Coordinate (State Plane)	Open Data Channel Type
Park Facility Name	Park Borough	Vehicle Type	Taxi Company Borough
Taxi Pick Up Location	Bridge Highway Name	Bridge Highway Direction	Road Ramp
Bridge Highway Segment	Latitude	Longitude	Location

Unique Key	Created Date	Closed Date	Agency	Agency Name	Complaint Type	Descriptor	Location Type	Incident Zip	Incident Address	Street Name
65903998	08/20/2025 02:22:00 PM	08/20/2025 07:00:00 PM	DEP	Department of Environmental Protect: Water System	Excessive Water In Basement (WEF)			11411	215-26 121 AVENUE	121 AVENUE
65349315	06/23/2025 10:38:00 PM	06/24/2025 12:50:00 AM	DEP	Department of Environmental Protect: Water System	Excessive Water In Basement (WEF)			11221	436 EVERGREEN AVENUE	EVERGREEN AVENUE
65151092	06/01/2025 12:32:00 PM	06/03/2025 01:50:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			11215	417 CARROLL STREET	CARROLL STREET
65012511	05/20/2025 01:26:00 PM	05/20/2025 01:40:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			10084	44 WATER STREET	WATER STREET
64854981	05/05/2025 11:56:00 AM	05/05/2025 02:25:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Odor (SAZ)			11201	59 PEARL STREET	PEARL STREET
65748414	01/12/2025 04:17:00 PM	01/12/2025 07:50:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			11360	215-24 23 AVENUE	23 AVENUE
65702004	01/09/2025 07:35:00 PM	01/09/2025 09:40:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			11360	215-20 23 ROAD	23 ROAD
65836640	01/05/2025 08:50:00 AM	01/05/2025 10:40:00 AM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			10461	1728 EASTCHESTER ROAD	EASTCHESTER ROAD
65924682	10/30/2024 11:59:00 AM	10/30/2024 02:09:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Odor (SAZ)			10302	241 CRYSTAL AVENUE	CRYSTAL AVENUE
62605459	10/18/2024 02:17:00 PM	10/18/2024 08:15:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Odor (SAZ)			10084	1 STATE STREET	STATE STREET
62211943	08/23/2024 04:50:00 PM	08/23/2024 06:35:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			10302	135 BRYSON AVENUE	BRYSON AVENUE
62056862	08/09/2024 01:05:00 PM	08/09/2024 02:15:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			10084	125 BROAD STREET	BROAD STREET
62057670	08/06/2024 09:44:00 AM	08/08/2024 12:00:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			10314	271 BIDWELL AVENUE	BIDWELL AVENUE
61376188	06/05/2024 11:53:00 AM	06/05/2024 01:23:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			11201	312 WATER STREET	WATER STREET
61270153	05/25/2024 04:08:00 PM	05/25/2024 05:50:00 PM	DEP	Department of Environmental Protect: Water System	Excessive Water In Basement (WEF)			11429	115-29 219 STREET	219 STREET
61073022	05/06/2024 11:17:00 PM	05/07/2024 09:25:00 AM	DEP	Department of Environmental Protect: Water System	LOW WATER PRESSURE - WLWP			11219	1249 42 STREET	42 STREET
61013228	04/29/2024 04:08:00 PM	04/29/2024 05:45:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Odor (SAZ)			11201	244 WATER STREET	WATER STREET
60683140	03/25/2024 03:38:00 PM	03/25/2024 06:23:00 PM	DEP	Department of Environmental Protect: Sewer	Manhole Overflow (Use Comments)			11360	215-44 23 ROAD	23 ROAD
59647995	12/05/2023 02:08:00 PM	12/05/2023 05:09:00 PM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			10314	299 BIDWELL AVENUE	BIDWELL AVENUE
59877008	11/09/2023 09:08:00 AM	11/09/2023 10:29:00 AM	DEP	Department of Environmental Protect: Sewer	Sewer Backup (Use Comments) (Sr)			10314	245 WILLARD AVENUE	WILLARD AVENUE

Weather Data Source:[Open-Meteo API](#)

Rows: A total of about 87,700 rows, with each row representing an hourly weather observation over a 10 year period.

Columns: 10 columns, including time, temperature, and precipitation measurements.

Columns of Interest or KPIs: Time, Precipitation, Rainfall, Windspeed.

5 Data Sets, For Each Borough - Data from October 1, 2015 to October 1, 2025

time	temperature_2m (°C)
precipitation (mm)	rain (mm)
cloudcover (%)	cloudcover_low (%)
cloudcover_mid (%)	cloudcover_high (%)
windspeed_10m (km/h)	winddirection_10m (°)

Sample - Queens Data

time	temperature_2m (°C)	precipitation (mm)	rain (mm)	cloud_cover (%)	cloud_cover_low (%)	cloud_cover_mid (%)	cloud_cover_high (%)	wind_speed_10m (km/h)	wind_direction_10m (°)
2015-10-03T00:00	0.4	0	0	58	5	9	53	4.3	175
2015-10-03T01:00	0.3	0	0	68	19	20	60	4.7	171
2015-10-03T02:00	0.9	0.3	0.3	88	25	37	83	3.4	162
2015-10-03T03:00	5.2	0.2	0.1	94	23	48	90	4	175
2015-10-03T04:00	6.6	0.3	0.2	96	18	82	93	4.8	228
2015-10-03T05:00	7.4	0.2	0.2	98	74	93	89	6.8	238
2015-10-03T06:00	8.2	0.1	0.1	99	76	93	60	10.3	245
2015-10-03T07:00	9.5	0.1	0.1	95	72	92	37	10.8	249
2015-10-03T08:00	8.4	0.1	0.1	94	69	81	35	9.9	251
2015-10-03T09:00	8.6	0.7	0.7	77	65	51	27	10.2	247
2015-10-03T10:00	8.9	0.2	0.2	60	48	24	20	7.4	247
2015-10-03T11:00	7.6	0.1	0.1	48	27	2	27	6.7	234
2015-10-03T12:00	7.3	0.1	0	51	22	15	41	5.8	240
2015-10-03T13:00	6.6	0.1	0.1	67	33	39	46	4.3	222
2015-10-03T14:00	5.7	0.1	0.1	66	37	60	1	3.4	212
2015-10-03T15:00	5.1	0.2	0.2	67	51	49	0	4	185
2015-10-03T16:00	4.8	0.4	0.4	59	50	36	0	5	180
2015-10-03T17:00	4.6	0	0	55	51	28	0	5.4	172
2015-10-03T18:00	3.7	0	0	69	65	29	1	5.1	172
2015-10-03T19:00	3	0	0	66	65	25	0	5.2	168
2015-10-03T20:00	2.8	0	0	52	49	16	0	5.4	160
2015-10-03T21:00	2.3	0	0	40	39	8	0	5.5	157
2015-10-03T22:00	1.9	0	0	29	22	13	0	4.4	171
2015-10-03T23:00	1.4	0	0	22	21	1	0	4.9	163

Potential Data Source 2:

[NYC Open Data – Street Flooding \(2010 to Present\)](#)

- Rows: A total of about 43 thousand rows where each row is a report.
- Columns: A total of 39 columns, including a unique identifier.
- Columns of Interest or KPIs: Complaint Type, Incident ZIP, Borough, Longitude, Latitude, Descriptor, Created Date, Closed Date, Resolution Description.

NYC Open Data Street Flooding Dataset Columns:

Unique Key	Created Date	Closed Date	Agency
Agency Name	Complaint Type	Descriptor	Location Type
Incident Zip	Incident Address	Street Name	Cross Street 1
Cross Street 2	Intersection Street 1	Intersection Street 2	Address Type
City	Landmark	Facility Type	Status
Due Date	Resolution Description	Resolution Action Updated Date	Community Board
Borough	X Coordinate (State Plane)	Y Coordinate (State Plane)	Park Facility Name
Park Borough	Vehicle Type	Taxi Company Borough	Taxi Pick Up Location
Bridge Highway Name	Bridge Highway Direction	Road Ramp	Bridge Highway Segment
Latitude	Longitude	Location	–

Other:

Potential APIs:

- [Socrata Open Data API \(SODA\)](#) – NYC Open Data API
- [National Centers for Environmental Information API \(NOAA\)](#) – Extreme Weather Events

Part II: Requirements Gathering & KPIs

Data Sources –

1. 311 Service Complaints (NYC Open Data)

Source: <https://data.cityofnewyork.us/Social-Services/311-Service-Requests-from-2010-to-Present>

Update Frequency: Daily

Time Period Being Used: Oct 2015 - Oct 2025

Columns Being Used: Unique Key, Created Date, Closed Date, Agency, Agency Name, Complaint Type, Descriptor, Location Type, Incident Zip, City, Status, Resolution Action Updated, Borough, Latitude, Longitude

2. Hourly Weather Data by Borough

Source: Open Meteo API

Update Frequency: Daily

Time Period Being Used: Oct 2015 - Oct 2025

Columns Being Used: Time, Temperature, Precipitation, Rain, Wind Speed

Key Performance Indicators (KPIs) –

311 Water Complaints KPIs

1. Daily total water-related complaints by borough.
2. Complaint count per 10,000 residents by borough/ZIP.
3. Average response time (Created → Closed) for water-related complaints.
4. Percent of complaints resolved within SLA (e.g., 7 days).
5. Severity breakdown by descriptor (flooding, sewer backup, leak, etc.).

Weather KPIs

1. Daily precipitation (rainfall + snowfall) per borough.
2. Number of extreme weather days (e.g., >2 inches rain, >6 inches snow).
3. Average precipitation per month by borough.

Integrated KPIs (311 + Weather) –

1. Correlation between daily precipitation and complaint counts by borough.
2. Average response times on extreme precipitation days vs. normal days.
3. Borough ranking of water complaints during heavy rain events.
4. Hotspot analysis: ZIP codes with highest complaint counts per inch of rain.
5. Service gap KPI: Difference in average resolution time during storm events between boroughs.

Part III: Dimensional Modeling

Kimball's Four Step Process

Step 1: Selecting the Business Processes – There are two main processes that are going to be modeled – 311 Complaints, and Weather Monitoring.

The 311 Complaints Process – This process will be modeled to capture incidents reported to NYC's 311 line, specifically complaints related to the water system, sewer system, and storm or flooding complaints. Each complaint includes information such as the complaint type, the location of the reported incident, the agency responsible for resolving the issue. This process provides insight into how city infrastructure is affected by water-related issues.

The Weather Monitoring Process – This process tracks hourly and daily weather data for each borough within the same time period as the 311 complaints. This process contains measures such as precipitation, rainfall, temperature, and wind speed.

Step 2: Declaring the Grain

For the 311 complaints, the grain is defined as one row for each individual complaint made that can be uniquely identified with the “Unique Key” provided by NYC Open Data. This allows measures such as resolution time and counts of incidents to be calculated accurately for each complaint. For the weather data, the grain is one row per hour for each borough, capturing weather conditions such as precipitation, rainfall, and windspeed.

Step 3: Identifying the Dimensions

There will be two dimensions shared between both fact tables – dim_date, a dimension containing time related attributes, and dim_location, which includes location attributes of the both the reported incidents and the weather.

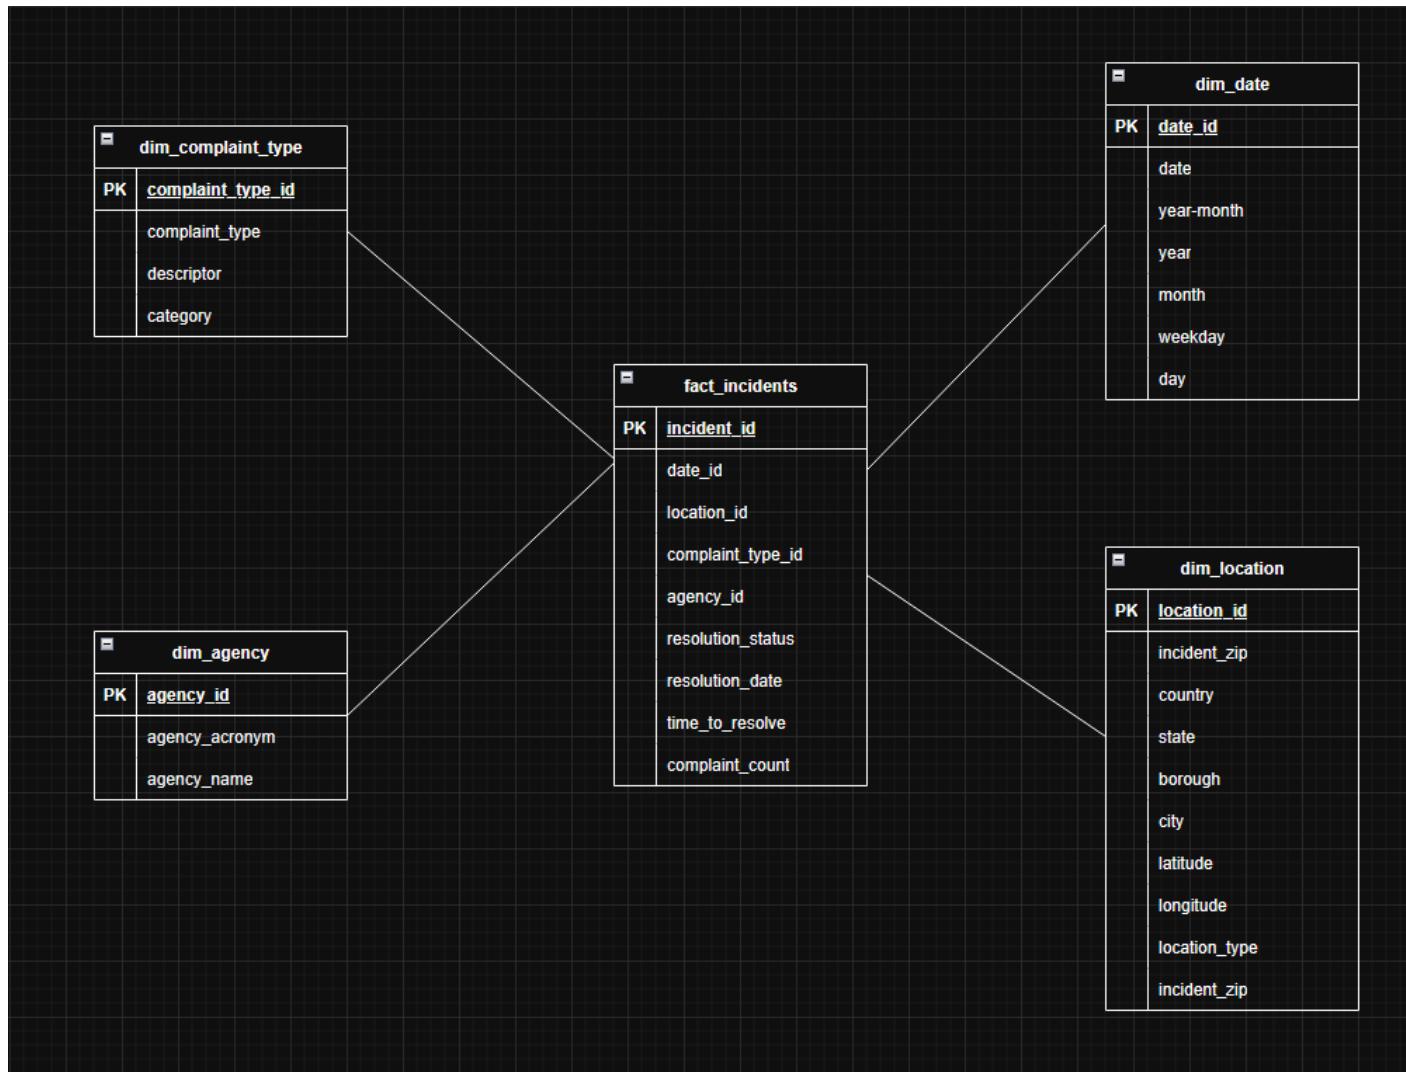
There will also be a dim_complaint_type, which includes the type of complaint made and description of the complaint. A dim_agency dimension will also be included to capture the agencies responsible for resolving the complaint.

Step 4: Identifying the Facts

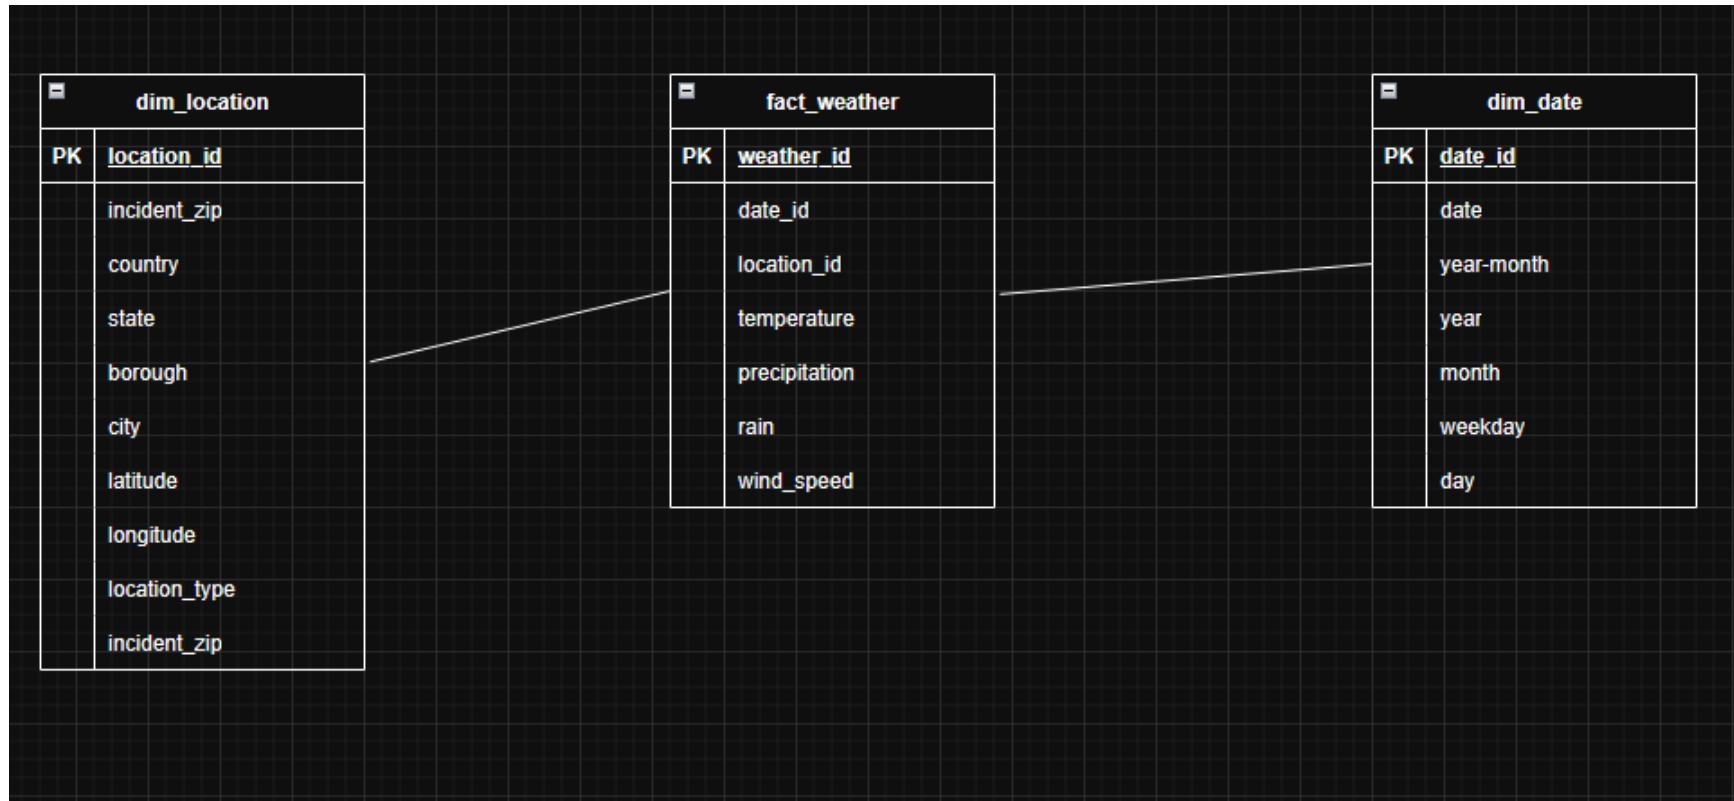
Fact_Incidents: Measures include incident_count (usually 1 per complaint) and Resolution_Time, calculated from Created Date to Resolution Action Updated Date. These measures allow for evaluation of complaint volume, response efficiency, and trends in water-related issues.

Fact_Weather: Measures include Temperature, Precipitation, Rain, Wind_Speed, Wind_Direction, and Cloud_Cover, enabling analysis of weather patterns and correlation with 311 complaints.

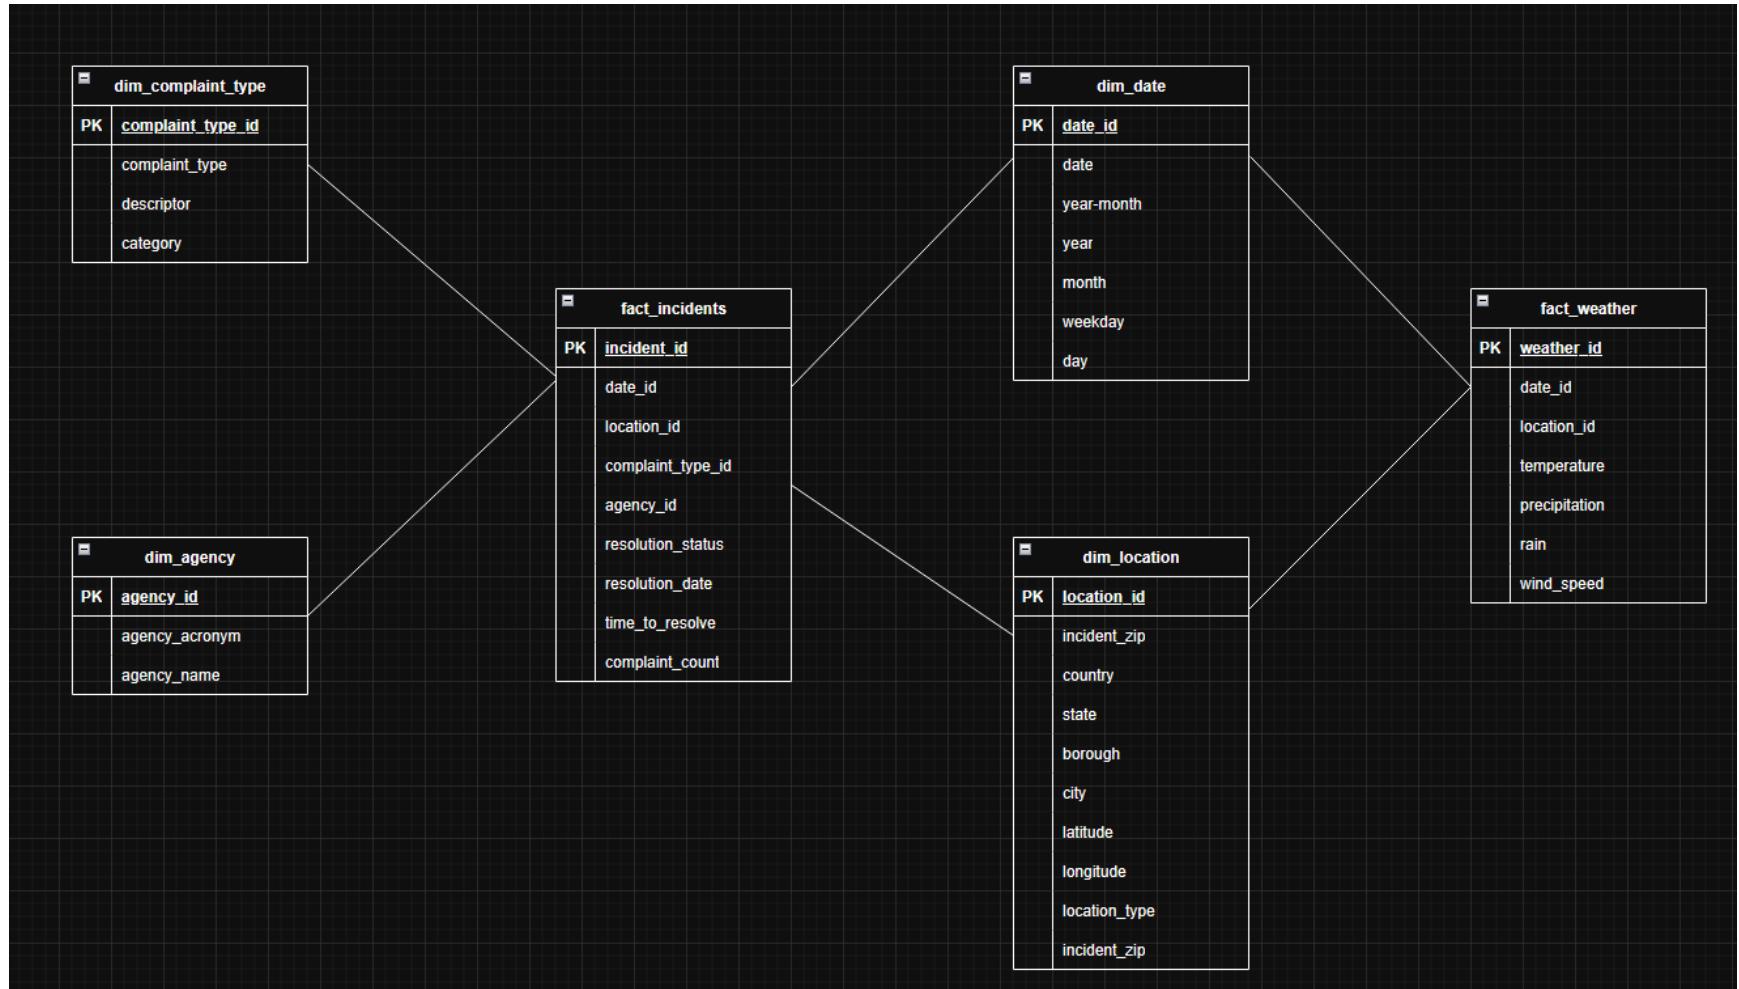
NYC Water-Related Complaints Star Schema



Weather Patterns Star Schema



NYC Water-Related Complaints & Weather Patterns Star Schema



Part IV: Technical Architecture

Target Data Warehouse: Google BigQuery

Hosting Environment: Google Cloud

Automation & Orchestration: Airflow

ETL in Python (Libraries & Frameworks):

- Extract: Urllib, PyArrow, Parquet, SODA, openmeteo-requests
- Transform: Pandas, Polars, DuckDB, Great Expectations
- Load: Google-Cloud-BigQuery

Data from the NYC 311 dataset will be extracted with Python by querying the Socrata Open Data API using URL queries with urllib to encode queries directly in the URL. The data is loaded in chunks and in parallel using Polars to handle large in-memory datasets efficiently. Each chunk is then converted into an Apache Arrow table using PyArrow for optimized constant reading/writing operations, and it will be written to a Parquet file, which provides efficient compression for storage. Weather data will be pulled using openmeteo-requests, processed with Pandas, and then saved as CSVs.

Data will be processed, profiled, and transformed using a combination of Polars for large in-memory preprocessing, and Pandas for exploratory data analysis. Visualization libraries such as Matplotlib and Seaborn (which integrate with Pandas) will be used to visualize missingness and cardinality. Great Expectations will be used for data validation and data quality assurance. DuckDB will be used as a staging database file to hold the table data for the final dimension and fact tables from the processed data.

Each DuckDB table will be downloaded as an individual Parquet file for efficient storage, and for easy integration with the target data warehouse. The data will then be loaded into the target data warehouse, Google BigQuery, with Google Cloud as the hosting environment for the data. The Google BigQuery Python API will be used for loading the data into the target data warehouse. Apache Airflow will be used for automation and orchestration of the ETL process once the initial pipeline is created.

Part V: ETL Programming

Our ETL pipeline was implemented in Python and organized into modular steps: (1) extracting raw 311 and weather data from the API's (2) Validating and transforming into clean and consistent schemas, (3) Staging curated output in efficient file formats(Parquet) and a staging engineer (DuckDB) and (4) Loading the final dimension and fact tables into Google BigQuery for analytics and dashboarding. This section documents the major ETL modules and highlights key custom code.

Link to Repo: <https://github.com/AGxData/nyc-311-weather-etl>

Look from left to right

Extract code for 311 dataset

```

1 # Importing Libraries needed for extraction
2 import polars as pl
3 import pyarrow as pa
4 import pyarrow.parquet as pq
5 import urllib.parse
6 import json
7 import os
8 from datetime import datetime
9 from concurrent.futures import ThreadPoolExecutor, as_completed
10 from frictionless import Resource, Schema, Field
11 from logger.etl_logger import ETLLogger
12 from pathlib import Path
13
14
15 # Setting Logging for extraction
16 extract_logger = ETLLogger("extract").get()
17 extract_logger.info("Starting 311 data extraction")
18
19 # Paths for project root, metadata, and full data
20 project_root = Path(__file__).resolve().parents[2]
21 main_parquet = project_root / "data" / "nyc_311_full_preprocessed.parquet"
22 metadata_folder = project_root / "metadata"
23
24 # Columns for extraction
25 columns = [
26     "unique_key", "created_date", "closed_date", "agency", "agency_name",
27     "complaint_type", "descriptor", "location_type", "incident_zip",
28     "city", "status", "resolution_action_updated_date", "borough",
29     "latitude", "longitude"
30 ]
31
32 # columns needed for overwriting
33 slowly_changing_dimensions = [
34     "created_date",
35     "closed_date",
36     "resolution_action_updated_date"
37 ]
38
39 # Frictionless schema for validation
40 schema = Schema(fields = [
41     Field(name = "unique_key", type = "string"),
42     Field(name = "created_date", type = "datetime"),
43     Field(name = "closed_date", type = "datetime"),
44     Field(name = "agency", type = "string"),
45     Field(name = "agency_name", type = "string"),
46     Field(name = "complaint_type", type = "string"),
47     Field(name = "descriptor", type = "string"),
48     Field(name = "location_type", type = "string"),
49     Field(name = "incident_zip", type = "string"),
50     Field(name = "city", type = "string"),
51     Field(name = "status", type = "string"),
52     Field(name = "resolution_action_updated_date", type = "datetime"),
53     Field(name = "borough", type = "string"),
54     Field(name = "latitude", type = "number"),
55     Field(name = "longitude", type = "number"),
56 ])
57
58 # Settings for extraction
59 base_url = r"https://data.cityofnewyork.us/resource/erm2-nwe9.csv"
60 chunk_size = 100_000
61 max_workers = 4
62
63
64
65 # Frictionless schema for validation upon extraction
66 schema = Schema(fields = [
67     Field(name = "unique_key", type = "string"),
68     Field(name = "created_date", type = "datetime"),
69     Field(name = "closed_date", type = "datetime"),
70     Field(name = "agency", type = "string"),
71     Field(name = "agency_name", type = "string"),
72     Field(name = "complaint_type", type = "string"),
73     Field(name = "descriptor", type = "string"),
74     Field(name = "location_type", type = "string"),
75     Field(name = "incident_zip", type = "string"),
76     Field(name = "city", type = "string"),

```

--- This section sets up the ETL process by importing the libraries used for data processing, API calls, and validation. We

also define a schema that lists the columns we expect from the 311 dataset and their data types. This schema is later used to check that the data we extract matches the expected structure. Defining the schema early helps catch data issues before the data is saved or merged.

```

77     Field(name = "status", type = "string"),
78     Field(name = "resolution_action_updated_date", type = "datetime"),
79     Field(name = "borough", type = "string"),
80     Field(name = "latitude", type = "number"),
81     Field(name = "longitude", type = "number"),
82   ])
83
84
85 # Function for downloading by chunks from Socrata API via URL (Faster i/o)
86 def download_chunk(offset, latest_date):
87     sql = f"""
88         SELECT {', '.join(columns)}
89         WHERE created_date > '{latest_date}'
90         LIMIT {chunk_size} OFFSET {offset}
91     """
92
93     encoded_query = urllib.parse.quote(sql, safe='')
94     url = f'{base_url}?${query}={encoded_query}'
95
96     try:
97         df_chunk = pl.read_csv(url, columns=columns, dtypes={"incident_zip": pl.Utf8})
98         if df_chunk.height == 0:
99             return None
100        return df_chunk
101    except Exception as e:
102        extract_logger.error(f"Error at offset {offset}: {e}")
103        return None
104
105
106 def extract_311():
107     # Determining latest date in metadata extraction files
108     metadata_folder.mkdir(parents = True, exist_ok = True)
109     metadata_file = metadata_folder / "last_date.json"
110
111     if metadata_file.exists():
112         with open(metadata_file) as f:
113             report_json = json.load(f)
114             latest_date = report_json.get("last_date")

```

--- The 311 dataset is very large, so we do not download it all at once. Instead, we pull the data in smaller chunks using LIMIT and OFFSET. Each chunk is read into memory using Polars, which is faster and more memory-efficient than standard tools for large datasets. Chunking allows us to safely extract millions of rows without running out of memory.

--- This step allows the ETL process to only pull new 311 records. The pipeline checks a metadata file that stores the most recent created_date

```

115         extract_logger.info(f"Using latest_date from metadata: {latest_date}")
116     else:
117         latest_date = "2025-09-25T01:44:42" # default start date
118         extract_logger.info(f"No metadata found, using default start date: {latest_date}")
119
120     offset = 0
121     all_chunks = []
122     finished = False
123
124     while not finished:
125         offsets = [offset + i * chunk_size for i in range(max_workers)]
126         results = []
127
128         with ThreadPoolExecutor(max_workers=max_workers) as executor:
129             futures = {executor.submit(download_chunk, o, latest_date): o for o in offsets}
130             for future in as_completed(futures):
131                 df_chunk = future.result()
132                 if df_chunk is not None and df_chunk.height > 0:
133                     results.append(df_chunk)
134
135         if not results: # Stopping if all chunks are empty
136             finished = True
137         else:
138             all_chunks.extend(results)
139             offset += max_workers * chunk_size
140
141     if not all_chunks:
142         extract_logger.info("No new 311 data to extract.")
143         return None
144
145     # Concatenating data
146     new_data = pl.concat(all_chunks, rechunk=True)
147
148     # Schema Validation
149     try:
150         resource = Resource(data = new_data.to_dicts(), schema = schema)
151         validation_report = resource.validate()
152         if validation_report.valid:

```

--- To speed up the extraction process, multiple chunks are downloaded at the same time using a thread pool. This means the program can request data from the API in parallel instead of waiting for each request to finish one at a time. This is because parallel downloads significantly reduce the total runtime of the ETL process.

processed in the last run. If the file exists, extraction starts from that date instead of reloading all historical data.

```

153         extract_logger.info("Extracted data matches schema.")
154     else:
155         extract_logger.warning("Schema validation failed. Check extracted data.")
156     except Exception as e:
157         extract_logger.error(f"Schema validation failed: {e}")
158
159     # Merging with main dataset with incremental updatation
160     if main_parquet.exists():
161         df_main = pl.read_parquet(main_parquet)
162
163     # Joining on SCD columns
164     df_update = new_data.select(["unique_key"] + slowly_changing_dimensions)
165     df_main = df_main.join(df_update, on="unique_key", how="left")
166
167     # Overwriting SCD columns if new value exists
168     for col in slowly_changing_dimensions:
169         right_col = f"{col}_right"
170         if right_col in df_main.columns:
171             df_main = df_main.with_columns(
172                 pl.when(pl.col(right_col).is_not_null())
173                     .then(pl.col(right_col))
174                     .otherwise(pl.col(col))
175                     .alias(col)
176             ).drop(right_col)
177
178     # Adding new rows that do not exist
179     new_rows = new_data.join(df_main, on="unique_key", how="anti")
180     if new_rows.height > 0:
181         df_main = pl.concat([df_main, new_rows], rechunk=True)
182
183     combined = df_main
184 else:
185     combined = new_data
186
187     # Updating metadata files
188     last_date = new_data.select(pl.col("created_date").max()).item()
189     with open(metadata_file, "w") as f:
190         ison.dump({"last_date": last_date}, f)

```

--- This part of the code merges newly extracted data with the existing dataset. If a complaint already exists, certain fields such as dates are updated when newer values are available. Complaints that do not already exist are added as new rows. This keeps the dataset accurate and up to date without creating duplicates.

```

191     # Saving updated dataset
192     combined.write_parquet(main_parquet)
193     extract_logger.info(f"Extraction completed. Total records: {combined.height}")
194
195     return combined
196
197
198
199 # Entry point for the extraction function
200 if __name__ == "__main__":
201     extract_311()

```

--- The final cleaned dataset is saved as a Parquet file, which is efficient for storage and analytics. The metadata file is then updated with the latest processed date so future runs know where to resume. This makes the ETL process repeatable and efficient over time.

Extract code for Weather dataset

```

1 import polars as pl
2 from datetime import datetime, timedelta
3 from concurrent.futures import ThreadPoolExecutor, as_completed
4 import requests
5 import json
6 from pathlib import Path
7 from logger.etl_logger import ETLLogger
8
9 # Logger settings
10 extract_logger = ETLLogger("extract_weather").get()
11 extract_logger.info("Starting weather data extraction")
12
13 # Settings for file path locations
14 project_root = Path(__file__).resolve().parents[2]
15 metadata_folder = project_root / "metadata"
16 metadata_folder.mkdir(parents=True, exist_ok=True)
17 metadata_file = metadata_folder / "weather_last_date.json"
18
19 if metadata_file.exists():
20     with open(metadata_file) as f:
21         metadata = json.load(f)
22     last_date = datetime.fromisoformat(metadata.get("last_date"))
23     extract_logger.info(f"Using last_date from metadata: {last_date.date()}")
24 else:
25     last_date = datetime(2010, 1, 1)
26     extract_logger.info(f"No metadata found. Starting from default date: {last_date.date()}")
27
28
29
30
31 # Settings for pulling data
32 chunk_days = 30
33 max_workers = 4
34 end_date = datetime(2025, 9, 25)
35
36
37
38

```

```

77     current_start = current_end + timedelta(days = 1)
78
79
80
81
82 # Function to help pull lat/lon borough centroid data
83 def fetch_weather(borough, lat, lon, start, end):
84     params = {
85         "latitude": lat,
86         "longitude": lon,
87         "start_date": start.strftime("%Y-%m-%d"),
88         "end_date": end.strftime("%Y-%m-%d"),
89         "daily": ",".join(variables),
90         "timezone": "America/New_York"
91     }
92     try:
93         resp = requests.get(base_url, params=params, timeout=60)
94         resp.raise_for_status()
95         data = resp.json()
96         if "daily" not in data:
97             extract_logger.warning(f"No daily data for {borough} ({start.date()} to {end.date()})")
98             return None
99         df = pl.DataFrame(data["daily"])
100        df = df.with_columns([
101            pl.lit(borough).alias("borough"),
102            pl.lit(lat).alias("latitude"),
103            pl.lit(lon).alias("longitude")
104        ])
105        return df
106    except Exception as e:
107        extract_logger.error(f"Error fetching {borough} ({start.date()} to {end.date()}): {e}")
108        return None
109
110
111
112
113 # Main function for weather extraction
114 def extract_weather():

```

```

39 # Open-Meteo variables of interest
40 variables = [
41     "temperature_2m_max",
42     "temperature_2m_min",
43     "precipitation_sum",
44     "precipitation_hours",
45     "rain_sum",
46     "showers_sum",
47     "snowfall_sum",
48     "windspeed_10m_max",
49     "windgusts_10m_max"
50 ]
51
52 # Open-Meteo base URL
53 base_url = "https://archive-api.open-meteo.com/v1/archive"
54
55
56
57
58
59 # Centroid Lat/Lon coordinates for each borough
60 borough_coords = {
61     "Manhattan": (40.7831, -73.9712),
62     "Brooklyn": (40.6782, -73.9442),
63     "Queens": (40.7282, -73.7949),
64     "Bronx": (40.8448, -73.8648),
65     "Staten Island": (40.5795, -74.1502)
66 }
67
68
69
70
71 # Function to help pull data in chunks based on date ranges
72 def daterange_chunks(start, end, days_per_chunk = 30):
73     current_start = start
74     while current_start <= end:
75         current_end = min(current_start + timedelta(days = days_per_chunk-1), end)
76         yield current_start, current_end

```

```

115     current_max_date = last_date
116     all_chunks = []
117
118     for chunk_start, chunk_end in daterange_chunks(last_date + timedelta(days = 1), end_date, chunk_days):
119         extract_logger.info(f"Fetching data from {chunk_start.date()} to {chunk_end.date()}")
120         dfs = []
121
122         with ThreadPoolExecutor(max_workers=max_workers) as executor:
123             futures = (executor.submit(fetch_weather, b, *borough_coords[b], chunk_start, chunk_end) for b in borough_coords)
124             for future in as_completed(futures):
125                 result = future.result()
126                 if result is not None:
127                     dfs.append(result)
128
129         if dfs:
130             chunk_df = pl.concat(dfs, rechunk=True)
131             all_chunks.append(chunk_df)
132
133             max_chunk_date = max(chunk_df["time"].to_list())
134             if isinstance(max_chunk_date, str):
135                 max_chunk_date = datetime.fromisoformat(max_chunk_date)
136             if max_chunk_date > current_max_date:
137                 current_max_date = max_chunk_date
138
139     if not all_chunks:
140         extract_logger.info("No new weather data to extract.")
141         return None
142
143     combined_df = pl.concat(all_chunks, rechunk=True)
144
145     # Updating metadata files
146     with open(metadata_file, "w") as f:
147         json.dump({"last_date": current_max_date.isoformat()}, f)
148     extract_logger.info(f"Weather extraction complete. last_date updated to {current_max_date.date()}")
149
150
151
152

```

```
154  
155  
156 # Entry point for weather extraction function  
157 if __name__ == "__main__":  
158     weather_df = extract_weather()
```

Transform code for 311 dataset

```

156     pl.col(column)
157     .map_elements(lambda x: fuzzy_map.get(x, x))
158     .alias(column)
159   )
160 transform_logger.info(f"Applied mapping on {column} with fuzzy matching: {use_fuzzy}")
161 return df
162
163
164
165 def transform_311(df: pl.DataFrame, mappings: dict) -> pl.DataFrame:
166   df = data_type_transformer(df)
167
168   df = clean_strings_before_mapping(df,
169     "complaint_type", "location_type", "city", "borough", "agency_name", "descriptor"
170   )
171
172   df = filter_relevant_complaints(df, mappings["relevant_complaints"])
173
174   df = dedupe(df)
175
176   mapping_columns = {
177     "complaint_type": "complaint_mapping",
178     "complaint_category": "complaint_categories",
179     "agency": "agency_mapping",
180     "agency_name": "agency_name_mapping",
181     "city": "city_mapping",
182     "borough": "borough_mapping",
183     "location_type": "location_type_mapping"
184   }
185   for col, mapping_name in mapping_columns.items():
186     df = apply_mapping(df, col, mappings.get(mapping_name, {}))
187
188   df = clean_zip_codes(df)
189
190   df = title_casing(df)
191
192   str_columns = [col for col, dtype in zip(df.columns, df.dtypes) if dtype == pl.Utf8]
193   df = df.with_columns([
194     pl.when(pl.col(col) == "missing")
195       .then(None)
196       .otherwise(pl.col(col))
197       .alias(col)
198     for col in str_columns
199   ])
200
201   return df
202
203
204 # Function entry point for the main 311 transformation function
205 if __name__ == "__main__":
206   transform_311()

```

Transform code for Weather dataset

```

1  # transform/transform_311_weather.py
2  import polars as pl
3
4  def transform_combined(cases: pl.DataFrame, weather: pl.DataFrame) -> dict:
5      # Ensuring data types
6      cases = cases.with_columns([
7          pl.col("unique_key").cast(pl.Int64),
8          pl.col("created_date").cast(pl.Date),
9          pl.col("closed_date").cast(pl.Date),
10         pl.col("resolution_action_updated_date").cast(pl.Date),
11         pl.col("agency").cast(pl.Utf8),
12         pl.col("agency_name").cast(pl.Utf8),
13         pl.col("complaint_type").cast(pl.Utf8),
14         pl.col("descriptor").cast(pl.Utf8),
15         pl.col("location_type").cast(pl.Utf8),
16         pl.col("incident_zip").cast(pl.Utf8),
17         pl.col("city").cast(pl.Utf8),
18         pl.col("status").cast(pl.Utf8),
19         pl.col("borough").cast(pl.Utf8),
20         pl.col("latitude").cast(pl.Float64),
21         pl.col("longitude").cast(pl.Float64),
22         pl.col("complaint_category").cast(pl.Utf8),
23     ])
24
25     cases = cases.with_columns([
26         pl.when(pl.col("closed_date") < pl.col("created_date"))
27             .then(None)
28             .otherwise(pl.col("closed_date"))
29             .alias("closed_date")
30     ])
31
32     weather = weather.with_columns([
33         pl.col("time").str.strptime(pl.Date, "%Y-%m-%d").alias("date"),
34         pl.col("temperature_2m_max").cast(pl.Float64),
35         pl.col("temperature_2m_min").cast(pl.Float64),
36         pl.col("precipitation_sum").cast(pl.Float64),
37         pl.col("precipitation_hours").cast(pl.Int64),
38         pl.col("rain_sum").cast(pl.Float64),
39         pl.col("showers_sum").cast(pl.Float64),
40     ])
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76

```

```

39     pl.col("snowfall_sum").cast(pl.Float64),
40     pl.col("windspeed_10m_max").cast(pl.Float64),
41     pl.col("windgusts_10m_max").cast(pl.Float64),
42     pl.col("borough").cast(pl.Utf8),
43     pl.col("latitude").cast(pl.Float64),
44     pl.col("longitude").cast(pl.Float64),
45 ])
46
47 # dim_date
48 dim_date = pl.concat([
49     cases.select(["created_date", "closed_date"]).melt().select("value").rename({"value": "date"}),
50     weather.select(["time"]).rename({"time": "date"})
51 ]).unique().with_columns([
52     pl.col("date").cast(pl.Date),
53     pl.arange(1, pl.count() + 1).alias("date_id")
54 ]).sort("date")
55
56 dim_date = dim_date.with_columns([
57     pl.col("date").dt.day().alias("day"),
58     pl.col("date").dt.month().alias("month"),
59     pl.col("date").dt.year().alias("year"),
60     pl.col("date").dt.weekday().alias("weekday"),
61     pl.col("date").dt.strftime("%A").alias("weekday_name")
62 ])
63
64 # dim_borough
65 dim_borough = pl.DataFrame({
66     "borough_id": [1, 2, 3, 4, 5],
67     "borough_name": ["Manhattan", "Brooklyn", "Queens", "Bronx", "Staten Island"]
68 })
69
70 # dim_location
71 weather_loc = weather.select([
72     pl.col("borough"),
73     pl.col("latitude"),
74     pl.col("longitude"),
75     pl.lit(None).cast(pl.Utf8).alias("city"),
76     pl.lit(None).cast(pl.Utf8).alias("location_type"),

```

```

77     pl.lit(None).cast(pl.Utf8).alias("incident_zip"),
78
79     dim_location = pl.concat([
80         cases.select(["borough", "latitude", "longitude", "city", "location_type", "incident_zip"]),
81         weather_loc
82     ]).unique().with_columns([
83         pl.arange(1, pl.count() + 1).alias("location_id")
84     ])
85
86     dim_location = dim_location.select([
87         "location_id", "incident_zip", "borough", "city", "location_type", "latitude", "longitude"
88     ])
89
90     # dim_agency
91     dim_agency = cases.select(["agency", "agency_name"]).unique().with_columns([
92         pl.arange(1, pl.count() + 1).alias("agency_id")
93     ])
94
95     # dim_complaint_type
96     dim_complaint_type = cases.select(["complaint_type", "descriptor", "complaint_category"]).unique().with_columns([
97         pl.arange(1, pl.count() + 1).alias("complaint_type_id")
98     ])
99
100    # fact_incidents
101    fact_incidents = cases.join(
102        dim_date.rename({"date": "created_date", "date_id": "created_date_id"}),
103        on="created_date", how="left"
104    ).join(
105        dim_date.rename({"date": "closed_date", "date_id": "closed_date_id"}),
106        on="closed_date", how="left"
107    ).join(
108        dim_location,
109        on="borough", "latitude", "longitude"], how="left"
110    ).join(
111        dim_borough, on="borough", how="left"
112    ).join(
113        dim_agency, on="agency", how="left"
114    ).join(

```

```

153     # Filtering cols
154     fact_weather = fact_weather.select([
155         "date_id", "borough_id", "temperature_max", "temperature_min", "precipitation_total",
156         "precipitation_hours", "rain_total", "showers_total", "snowfall_total", "windspeed_max",
157         "windgust_max", "rain_flag", "showers_flag", "snow_flag", "high_wind_flag"
158     ])
159
160     # fact_daily_summary
161     daily_incidents = fact_incidents.group_by(["created_date_id", "borough_id"]).agg([
162         pl.count("incident_id").alias("total_incidents"),
163         (pl.col("is_resolved_same_day").mean() * 100).alias("percent_resolved_same_day")
164     ])
165
166     daily_weather = fact_weather.groupby(["date_id", "borough_id"]).agg([
167         pl.col("temperature_max").mean().alias("temperature_max"),
168         pl.col("temperature_min").mean().alias("temperature_min"),
169         ((pl.col("temperature_max") + pl.col("temperature_min")) / 2).mean().alias("temperature_avg"),
170         pl.col("precipitation_total").sum().alias("precipitation_total"),
171         (pl.when(pl.col("precipitation_total") > 0).then(pl.col("precipitation_total") / 24).otherwise(0)).alias("precipitation_per_hour"),
172         pl.col("rain_total").sum().alias("rain_total"),
173         pl.col("shower_total").sum().alias("showers_total"),
174         pl.col("snowfall_total").sum().alias("snowfall_total"),
175         pl.col("windspeed_max").max().alias("windspeed_max"),
176         pl.col("windgust_max").max().alias("windgust_max"),
177         pl.col("rain_flag").max().alias("rain_flag"),
178         pl.col("showers_flag").max().alias("showers_flag"),
179         pl.col("snow_flag").max().alias("snow_flag"),
180         pl.col("high_wind_flag").max().alias("high_wind_flag")
181     ])
182
183     fact_daily_summary = daily_incidents.join(
184         daily_weather, left_on="created_date_id", "borough_id", right_on="date_id", "borough_id", how="left"
185     )
186
187     # Renaming columns to match BigQuery schema
188     dim_date = dim_date.rename({
189         "date_id": "date_id", "date": "date", "day": "day", "month": "month",
190         "year": "year", "weekday": "weekday", "weekday_name": "weekday_name"
191     })

```

```

229     })
230
231     # Return all tables in a dict
232     return {
233         "dim_date": dim_date,
234         "dim_borough": dim_borough,
235         "dim_location": dim_location,
236         "dim_agency": dim_agency,
237         "dim_complaint_type": dim_complaint_type,
238         "fact_incidents": fact_incidents,
239         "fact_weather": fact_weather,
240         "fact_daily_summary": fact_daily_summary
241     }

```

```

115     dim_complaint_type,
116     one["complaint_type", "descriptor", "complaint_category"], how="left"
117 ).with_columns([
118     ((pl.col("closed_date") - pl.col("created_date")).alias("time_to_resolve_interval"),
119      ((pl.col("closed_date") == pl.col("created_date")).cast(pl.int64)).alias("is_resolved_same_day"),
120      pl.lit(1).alias("complaint_count"))
121 ])
122
123     # Dropping unnecessary columns
124     drop_cols_incidents = [
125         "agency_name", "complaint_type", "descriptor", "location_type",
126         "incident_zip", "city", "status", "resolution_action_updated_date",
127         "borough", "latitude", "longitude"
128     ]
129     fact_incidents = fact_incidents.drop([col for col in drop_cols_incidents if col in fact_incidents.columns])
130
131     # fact_weather
132     weather = weather.with_columns((pl.col("time").cast(pl.Date).alias("date")))
133
134     fact_weather = weather.rename({
135         "temperature_2m_max": "temperature_max",
136         "temperature_2m_min": "temperature_min",
137         "precipitation_sum": "precipitation_total",
138         "rain_sum": "rain_total",
139         "showers_sum": "showers_total",
140         "snowfall_sum": "snowfall_total",
141         "windspeed_10m_max": "windspeed_max",
142         "windgusts_10m_max": "windgust_max"
143     }).join(
144         dim_date.rename({"date": "date", "date_id": "date_id"}), on="date", how="left"
145     ).join(
146         dim_borough, on="borough", how="left"
147     ).with_columns([
148         (pl.col("rain_total") > 0).cast(pl.int64).alias("rain_flag"),
149         (pl.col("showers_total") > 0).cast(pl.int64).alias("showers_flag"),
150         (pl.col("snowfall_total") > 0).cast(pl.int64).alias("snow_flag"),
151         ((pl.col("windspeed_max") > 15) | (pl.col("windgust_max") > 20)).cast(pl.int64).alias("high_wind_flag")
152     ])
153

```

```

191     })
192     dim_borough = dim_borough.rename({"borough_id": "borough_id", "borough_name": "borough_name"})
193     dim_location = dim_location.rename({
194         "location_id": "location_id", "incident_zip": "incident_zip", "borough": "borough",
195         "city": "city", "location_type": "location_type", "latitude": "latitude", "longitude": "longitude"
196     })
197     dim_agency = dim_agency.rename({
198         "agency_id": "agency_id", "agency": "agency", "agency_name": "agency_name"
199     })
200     dim_complaint_type = dim_complaint_type.rename({
201         "complaint_type_id": "complaint_type_id", "complaint_type": "complaint_type",
202         "complaint_category": "complaint_category", "descriptor": "complaint_descriptor"
203     })
204     fact_incidents = fact_incidents.rename({
205         "unique_key": "incident_id", "created_date_id": "created_date_id", "closed_date_id": "closed_date_id",
206         "agency_id": "agency_id", "complaint_type_id": "complaint_type_id", "date_id": "date_id",
207         "location_id": "location_id", "resolution_status": "resolution_status",
208         "time_to_resolve_interval": "time_to_resolve_interval",
209         "is_resolved_same_day": "is_resolved_same_day",
210         "complaint_count": "complaint_count"
211     })
212     fact_weather = fact_weather.rename({
213         "date_id": "date_id", "borough_id": "borough_id", "temperature_max": "temperature_max",
214         "temperature_min": "temperature_min", "precipitation_total": "precipitation_total",
215         "precipitation_hours": "precipitation_hours", "rain_total": "rain_total",
216         "showers_total": "showers_total", "snowfall_total": "snowfall_total", "windspeed_max": "windspeed_max",
217         "windgust_max": "windgust_max", "rain_flag": "rain_flag", "showers_flag": "showers_flag",
218         "snow_flag": "snow_flag", "high_wind_flag": "high_wind_flag"
219     })
220     fact_daily_summary = fact_daily_summary.rename({
221         "created_date_id": "date_id", "borough_id": "borough_id", "total_incidents": "total_incidents",
222         "percent_resolved_same_day": "percent_resolved_same_day", "temperature_avg": "temperature_avg",
223         "temperature_max": "temperature_max", "temperature_min": "temperature_min",
224         "precipitation_total": "precipitation_total", "precipitation_per_hour": "precipitation_per_hour",
225         "rain_total": "rain_total", "showers_total": "showers_total", "snowfall_total": "snowfall_total",
226         "windspeed_max": "windspeed_max", "windgust_max": "windgust_max",
227         "rain_flag": "rain_flag", "showers_flag": "showers_flag", "snow_flag": "snow_flag",
228         "high_wind_flag": "high_wind_flag"
229     })

```

Load code for data

```

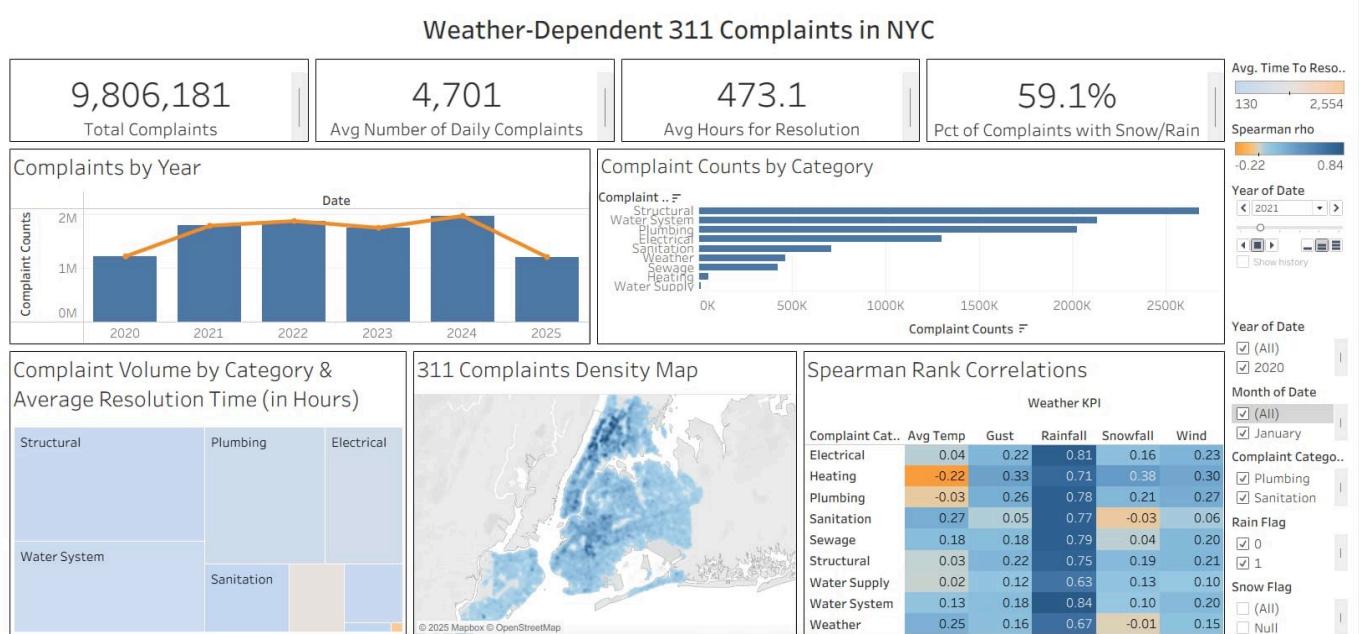
1 1|from google.cloud import bigquery
2 2|from google.oauth2 import service_account
3 3|import polars as pl
4 4|from pathlib import Path
5 5|from logger.etl_logger import ETLLogger
6 6|
7 7|# Initializing Logger
8 8|load_logger = ETLLogger("load").get()
9 9|load_logger.info("Starting data loading")
10 10|
11 11|# BigQuery client setup
12 12|script_dir = Path(__file__).parent
13 13|credentials_path = script_dir.parent / "credentials" / "bigquery_nyc_weather_etl_credentials.json"
14 14|credentials = service_account.Credentials.from_service_account_file(credentials_path)
15 15|client = bigquery.Client(project="nyc-311-weather-etl", credentials=credentials)
16 16|dataset_id = "nyc_311_weather"
17 17|
18 18|def load_to_bigquery(df_dict, chunk_size = 10_000):
19 19|    for table_name, df in df_dict.items():
20 20|        if df is None or df.height == 0:
21 21|            load_logger.info(f"No data for table {table_name}, skipping...")
22 22|            continue
23 23|        load_logger.info(f"Starting load for table {table_name} ({df.height} rows)")
24 24|        table_ref = f"{client.project}.{dataset_id}.{table_name}"
25 25|        # Check if table exists
26 26|        try:
27 27|            table = client.get_table(table_ref)
28 28|            table_exists = True
29 29|        except Exception:
30 30|            table_exists = False
31 31|            load_logger.info(f"Table {table_name} does not exist. Will create new table automatically.")
32 32|        # Convert Polars to pandas for BigQuery
33 33|        df_pd = df.to_pandas()
34 34|
35 35|        if table_name.startswith("dim_") and table_exists:
36 36|            # For dimension tables, remove duplicates based on primary key
37 37|            pk_field = table.schema[0].name
38 38|            # Load existing primary keys
39 39|existing_df = client.list_rows(table).to_dataframe()
40 40|df_pd = df_pd.merge(existing_df[[pk_field]], on=pk_field, how="left", indicator=True)
41 41|df_pd = df_pd[df_pd["_merge"] == "left_only"].drop(columns=["_merge"])
42 42|load_logger.info(f"{len(df_pd)} new rows detected for {table_name}")
43 43|
44 44|if len(df_pd) == 0:
45 45|    load_logger.info(f"No new rows to load for {table_name}")
46 46|    continue
47 47|
48 48|# Loading in chunks
49 49|for start in range(0, len(df_pd), chunk_size):
50 50|    end = start + chunk_size
51 51|    chunk_df = df_pd.iloc[start:end]
52 52|
53 53|try:
54 54|    job = client.load_table_from_dataframe(
55 55|        chunk_df,
56 56|        table_ref,
57 57|        job_config=bigquery.LoadJobConfig(write_disposition="WRITE_APPEND")
58 58|    )
59 59|    job.result() # Wait for completion
60 60|    load_logger.info(f"Loaded rows {start} to {end} into {table_name}")
61 61|except Exception as e:
62 62|    load_logger.error(f"Error loading rows {start} to {end} into {table_name}: {e}")
63 63|
64 64|load_logger.info(f"Finished loading table {table_name}")
65 65|
66 66|if __name__ == "__main__":
67 67|    load_to_bigquery()

```

Part VI: Visuals and Dashboard

BI Application

Our Tableau dashboard uses data from the integrated data warehouse to analyze water-related 311 complaints and their relationship to weather conditions across New York City. Below, we describe four key visualizations in detail, followed by brief descriptions of the remaining dashboard elements.



Summary KPI Metrics

This visualization shows high-level summary metrics, including total water-related 311 complaints, average daily complaints, average resolution time in hours, and the percentage of complaints that occurred on days with rain or snow. These KPIs update based on selected filters and provide a quick overview of complaint volume, response efficiency, and weather-related impact.

Complaints by Year

This chart displays the total number of water-related 311 complaints by year using bars, with a line showing the overall trend. It allows us to compare complaint volume across years and identify periods where water-related issues increased or decreased.

Complaint Counts by Category

This horizontal bar chart shows the total number of complaints by category, such as Structural, Water System, Plumbing, and Sanitation. The categories are ordered by volume, making it easy to see which types of issues are reported most frequently.

Complaint Volume by Category and Average Resolution Time

This heatmap visualizes complaint volume by category, where the size of each block represents the number of complaints and the color reflects the average resolution time. It helps compare both workload and service efficiency across different types of water-related issues.

311 Complaints Density Map

This map shows the geographic density of water-related 311 complaints across New York City. Darker areas indicate higher concentrations of complaints, allowing us to identify neighborhoods where water-related infrastructure issues are more common.

Spearman Rank Correlations Between Weather and Complaints

This heatmap displays the Spearman rank correlation between complaint categories and weather variables such as rainfall, snowfall, wind and temperature. The stronger color shading indicates stronger relationships, helping us understand which complaint types are most influenced by weather conditions.

Part VI: Conclusion

For our group, we used multiple tools to complete the project. Our main form of communication was WhatsApp, which we used for instant messaging and to set up Zoom or Google Meet calls. Through these platforms, we decided on our project and which softwares to use.

- Open-Meteo API: We used this to gather data on weather in the 5 boroughs, to correspond to our 311 Complaint data
- draw.io: We used this to design our dimensional models
- python: We used this to design and complete our ETL Programming
- Tableau: We used this to create visualizations for our data

The most difficult step in this process was completing the ETL Programming. We experienced a lot of issues with loading the data via Python, which suffered from local hardware limits. The easiest step was designing the Dimensional Model because we had already decided on the KPIs and relevant dimensions.

The results of our business application show that the proposed benefits of the system can be realized. The dashboard indicates that water-related 311 complaints increase on days with rain or snow, and that certain complaint categories, such as structural issues, appear more frequently during periods of heavy rainfall. The analysis also shows that some complaint types not only occur more often but also take longer to resolve, suggesting that severe weather places additional strain on city infrastructure and services. The category level and correlation analyses further support this by highlighting clear relationships between precipitation levels and specific types of complaints.

Overall, this project demonstrates the value of combining 311 service request data with weather data in a centralized data warehouse. By integrating multiple data sources and using business intelligence tools, we were able to identify patterns that would not be visible using a single dataset. The results of this application can be used to better understand how weather-related events impact NYC residents and to prepare more targeted responses as rainfall is expected to increase. For example, when heavy rainfall is forecasted, additional attention and resources can be directed toward areas that frequently report structural complaints, allowing the city to focus on failing infrastructure and provide increased support where it is most needed.

Part VII: Meeting Log

Meeting Log			
Date/Time	Modality	Attendees	Objective
9/9 16:00-17:00	Zoom	Adrian, Neha, Adnan, Roland	Filter 311 Data Set
9/10 14:30-15:15	Whatsapp	Adrian, Neha, Adnan, Roland	Discuss Secondary Data Sets
9/12 18:00-23:00	Whatsapp	Adrian, Neha, Adnan, Roland	Finalize Milestone 1 Submission
9/20 13:30-14:30	Zoom	Adrian, Neha, Adnan, Roland	Discuss Secondary Data Sets, Pick out KPIs
9/26 14:00-16:30	Zoom	Adrian, Neha, Adnan, Roland	Data Mart Schema
10/3 18:00-23:00	Zoom	Adrian, Neha, Adnan, Roland	Finalize Milestone 2 Submission
10/21 18:00-20:00	Zoom	Adrian, Neha, Adnan, Roland	Discuss and Finalize Milestone 3 Submission
11/7 18:00-19:00	Whatsapp	Adrian, Neha, Adnan, Roland	Discuss ETL Tools and DBMS
11/14 17:00-18:00	Zoom	Adrian, Neha, Adnan, Roland	Discussed Technical Architecture Components
12/1 18:00-20:00	Whatsapp	Adrian, Neha, Adnan, Roland	Discussed ETL Programming

12/5 16:00-20:00	Zoom	Adrian, Neha, Adnan, Roland	Discussed ETL Programming
12/19 16:00-18:00	Zoom	Adrian, Neha, Adnan, Roland	Finalized Visuals and Dashboard on Tableau

References

Data Source - 311 Complaint Data

https://data.cityofnewyork.us/Social-Services/311-Service-Requests-from-2010-to-Present/erm2-nwe9/about_data

Data Source - Weather In 5 Boroughs <https://open-meteo.com/>

Technique - Designing Star Schema

<https://www.geeksforgeeks.org/data-analysis/designing-the-star-schema-in-data-warehousing/>

Technique - ETL Programming in Python <https://www.datacamp.com/courses/etl-and-elt-in-python>

Technique - Creating a Dashboard <https://youtu.be/vDgBCgxLWPY?si=5AW5bkl5S4PbOd5I>